



Transforming Streams

Advanced stream processing in Rust

Willem Vanhulle

EuroRust 2025

1980-01-01

1. Introduction	1
1.1. Me	2
1.2. Kinds of streams	3
1.3. Streams in Rust are not new	4
1.4. Why does Rust need special treatment?	5
1.5. Process TCP connections and collect long messages	6
1.6. Stream operators: declarative & composable	7
2. Rust's Stream trait	8
3. Using Streams	14
4. Example 1: $1 \rightarrow 1$ Operator	20
5. Example 2: $1 \rightarrow N$ Operator	44
6. Conclusion	51



Motivation

Processing data from moving vehicles

1. Vehicle generates multiple data streams
2. All streams converge to control system

1.1. Me

Lives in Ghent, Belgium:

- Studied mathematics, physics and computer science
- Biotech automation (fermentation)
- Distributed systems (trains)



Motivation

Processing data from moving vehicles

1. Vehicle generates multiple data streams
2. All streams converge to control system

1.1. Me

Lives in Ghent, Belgium:

- Studied mathematics, physics and computer science
- Biotech automation (fermentation)
- Distributed systems (trains)

Latest projects (github.com/wvhulle):



- SysGhent.be: social network for systems programmers in Ghent (Belgium)
- Clone-stream: lazy stream cloning library for Rust

Motivation

Processing data from moving vehicles

1. Vehicle generates multiple data streams
2. All streams converge to control system



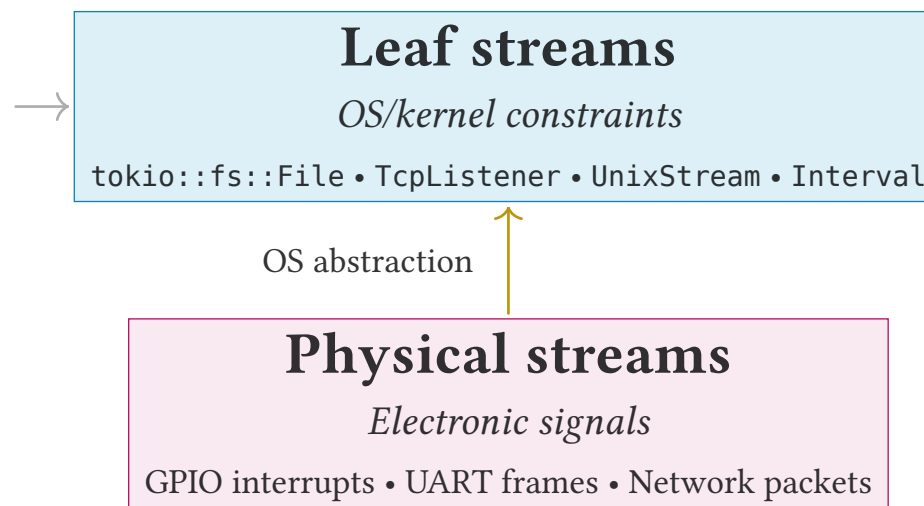
Physical streams

Electronic signals

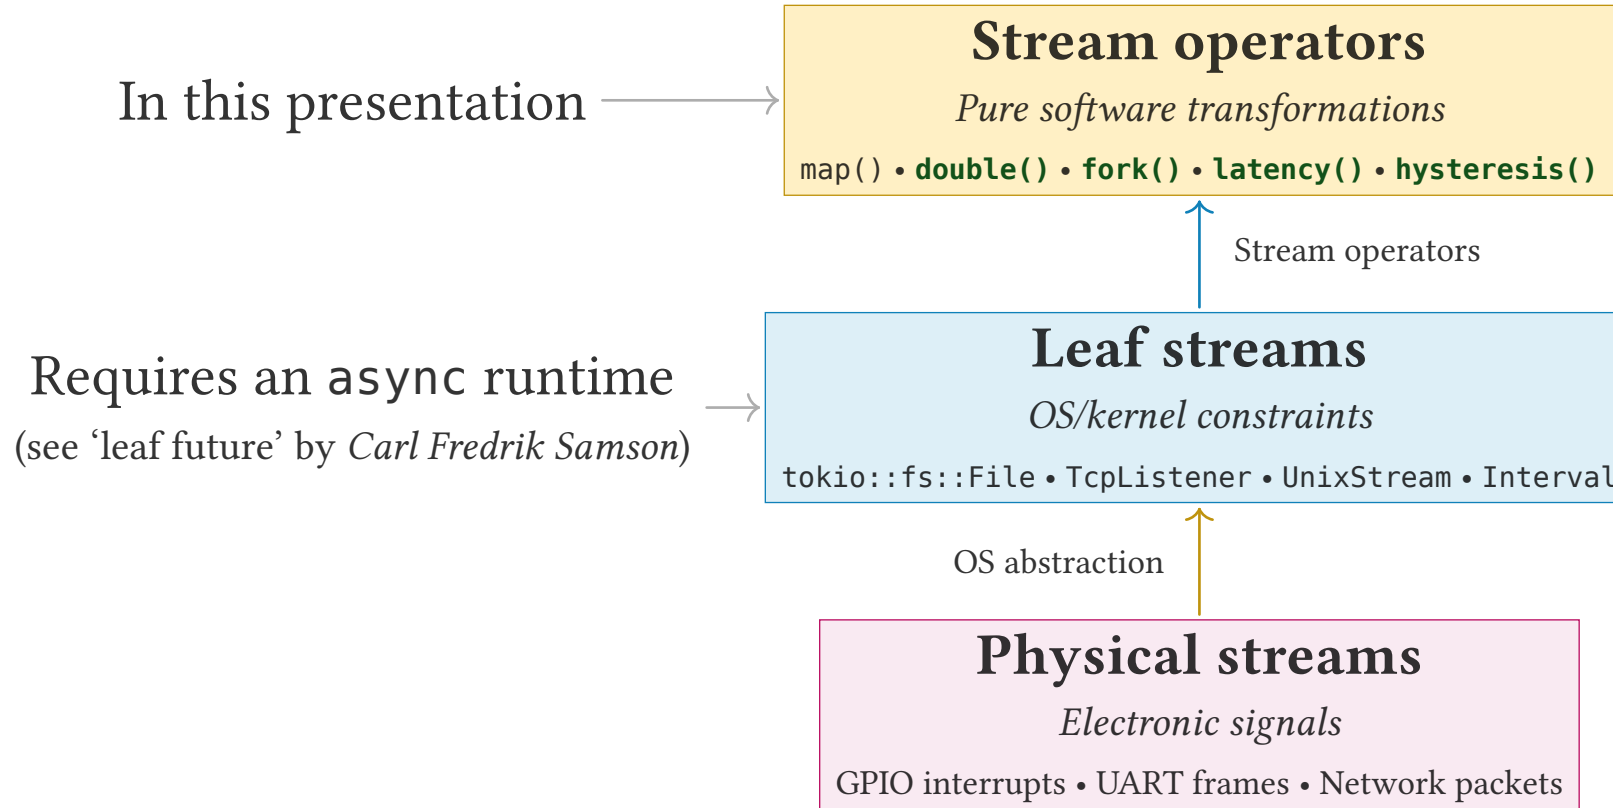
GPIO interrupts • UART frames • Network packets



Requires an async runtime
(see 'leaf future' by *Carl Fredrik Samson*)



1.2. Kinds of streams

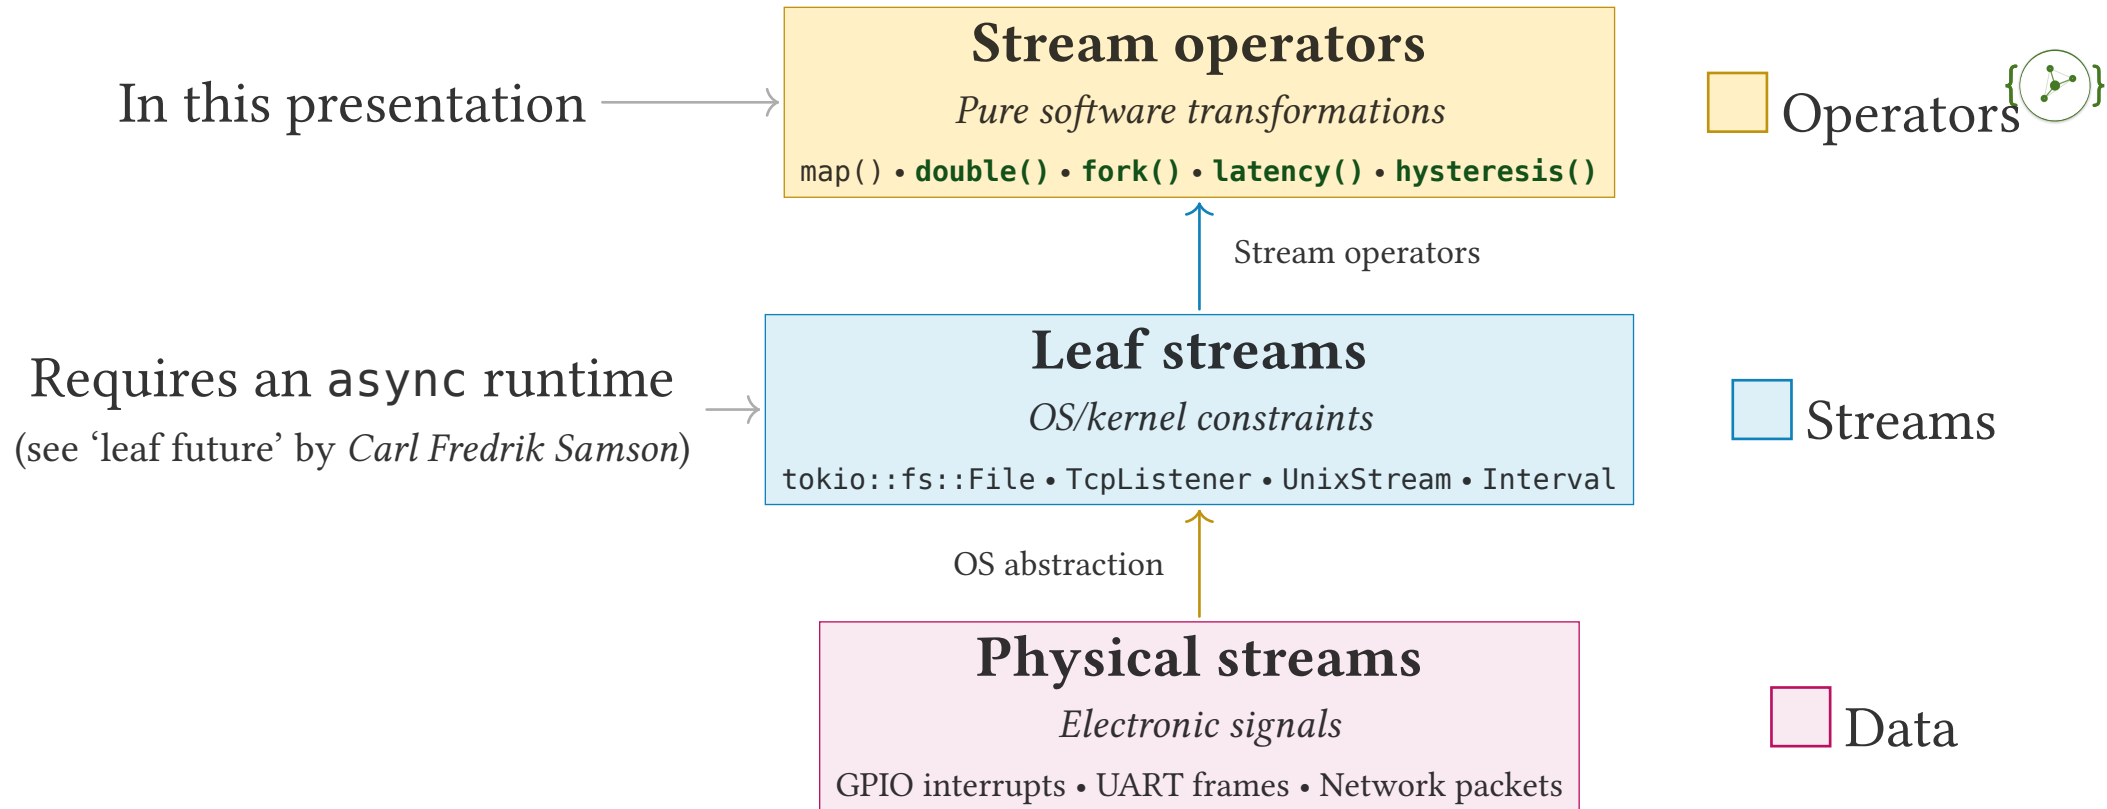


Operators 

Streams

Data

1.2. Kinds of streams



Hardware signals are abstracted by the OS

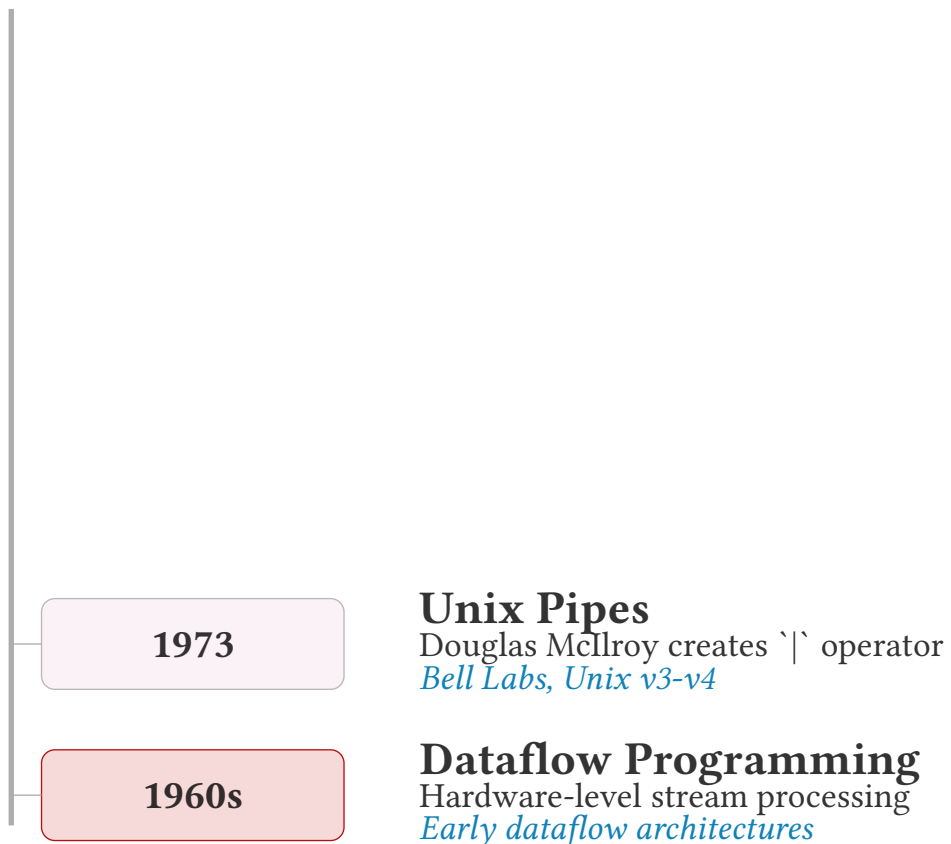
Software operators transform the streams




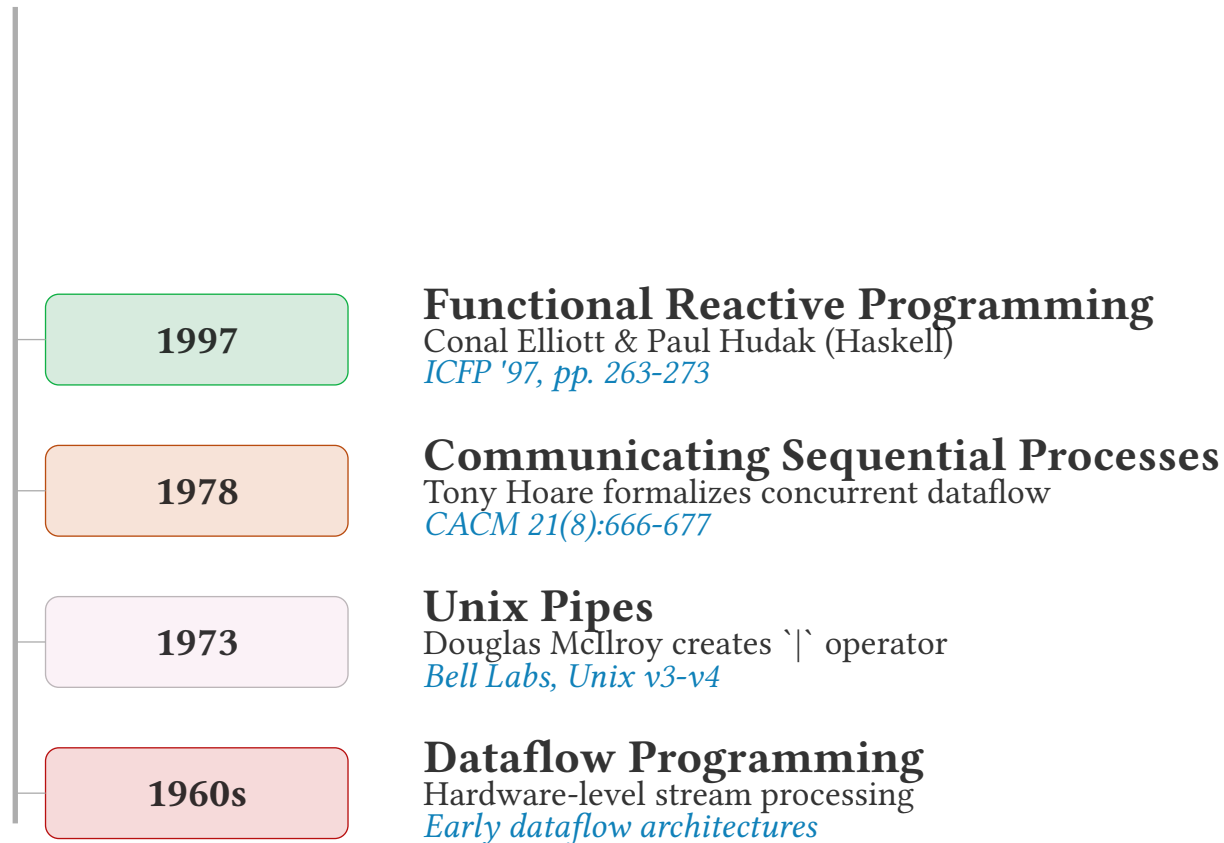



1960s

Dataflow Programming
Hardware-level stream processing
Early dataflow architectures













1.4. Why does Rust need special treatment?

Stream operators must wrap and own their input by value



vs



1.5. Process TCP connections and collect long messages



```
1  let mut results = Vec::new(); let mut count = 0;
2
3  while let Some(connection) = tcp_stream.next().await {
4      match connection {
5          Ok(stream) if should_process(&stream) => {
6              match process_stream(stream).await {
7                  Ok(msg) if msg.len() > 10 => {
8                      results.push(msg);
9                      count += 1;
10                     if count >= 5 { break; }
11                 }
12                 Ok(_) => continue,
13                 Err(_) => continue,
14             }
15         }
16         Ok(_) => continue,
17         Err(_) => continue,
18     }
19 }
```



Problems:

- Deeply nested
- Hard to read
- Cannot test pieces independently



1.6. Stream operators: declarative & composable

Same logic with stream operators:

```
1 let results: Vec<String> = tcp_stream
2   .filter_map(|conn| ready(conn.ok()))
3   .filter(|stream| ready(should_process(stream)))
4   .then(|stream| process_stream(stream))
5   .filter_map(|result| ready(result.ok()))
6   .filter(|msg| ready(msg.len() > 10))
7   .take(5)
8   .collect()
9   .await;
```

 Rust

Benefits:

- Each operation is isolated
- Testable
- Reusable

“Programs must be written **for people to read**”

1. Introduction	1
2. Rust's Stream trait	8
2.1. Moving from Iterator to Stream	9
2.2. The Stream trait: async iterator	10
2.3. Possible inconsistency	11
2.4. The meaning of Ready(None)	12
2.5. 'Fusing' Streams and Futures	13
3. Using Streams	14
4. Example 1: $1 \rightarrow 1$ Operator	20
5. Example 2: $1 \rightarrow N$ Operator	44
6. Conclusion	51

2.1. Moving from Iterator to Stream

✓ Always returns immediately

⚠ May be Pending

2. Rust's Stream trait

✓ Hides polling complexity



2.1. Moving from Iterator to Stream

✓ Always returns immediately

⚠ May be Pending

✓ Hides polling complexity



next() → Some(3)

next() → Some(1)

next() → None

next() → Some(2)

Iterator (sync)

2.1. Moving from Iterator to Stream

✓ Always returns immediately

⚠ May be Pending

✓ Hides polling complexity



next() → Some(3)

poll_next() → Ready(Some(2))

next() → Some(1)

poll_next() → Pending

next() → None

poll_next() → Ready(Some(1))

next() → Some(2)

poll_next() → Pending

Iterator (sync)

Stream (low-level)

2.1. Moving from Iterator to Stream

✓ Always returns immediately

⚠ May be Pending

✓ Hides polling complexity



next() → Some(3)

next() → Some(1)

next() → None

next() → Some(2)

Iterator (sync)

poll_next() → Ready(Some(2))

poll_next() → Pending

poll_next() → Ready(Some(1))

poll_next() → Pending

Stream (low-level)

next().await → Some(3)

next().await → Some(1)

next().await → None

next().await → Some(2)

Stream (high-level)



Actions



Data values



State




Stream

2.2. The Stream trait: async iterator

Like Future, but yields **multiple items** over time when polled:



```
1 trait Stream {  
2     type Item;  
3  
4     fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)  
5         -> Poll<Option<Self::Item>>;  
6 }
```



The Poll<Option<Item>> return type:

- Poll::Pending - not ready yet, try again later
- Poll::Ready(Some(item)) - here's the next item
- Poll::Ready(None) - stream is exhausted (no more items **right now**)

2.3. Possible inconsistency

```
1 trait Stream {  
2     type Item;  
3  
4     fn poll_next(self: Pin<&mut Self>, cx: &mut Context)  
5         -> Poll<Option<Self::Item>>  
6 }
```



Warning

What about Rust rule self needs to be Deref<Target=Self>?

`Pin<&mut Self>` only implements `Deref<Target=Self>` for `Self: Unpin`.

Problem? No, `Pin` is an exception in the compiler.



Regular Stream

“No items **right now**”

(Stream might yield more later)

Fused Stream

“No items **ever again**”

(Stream is permanently done)

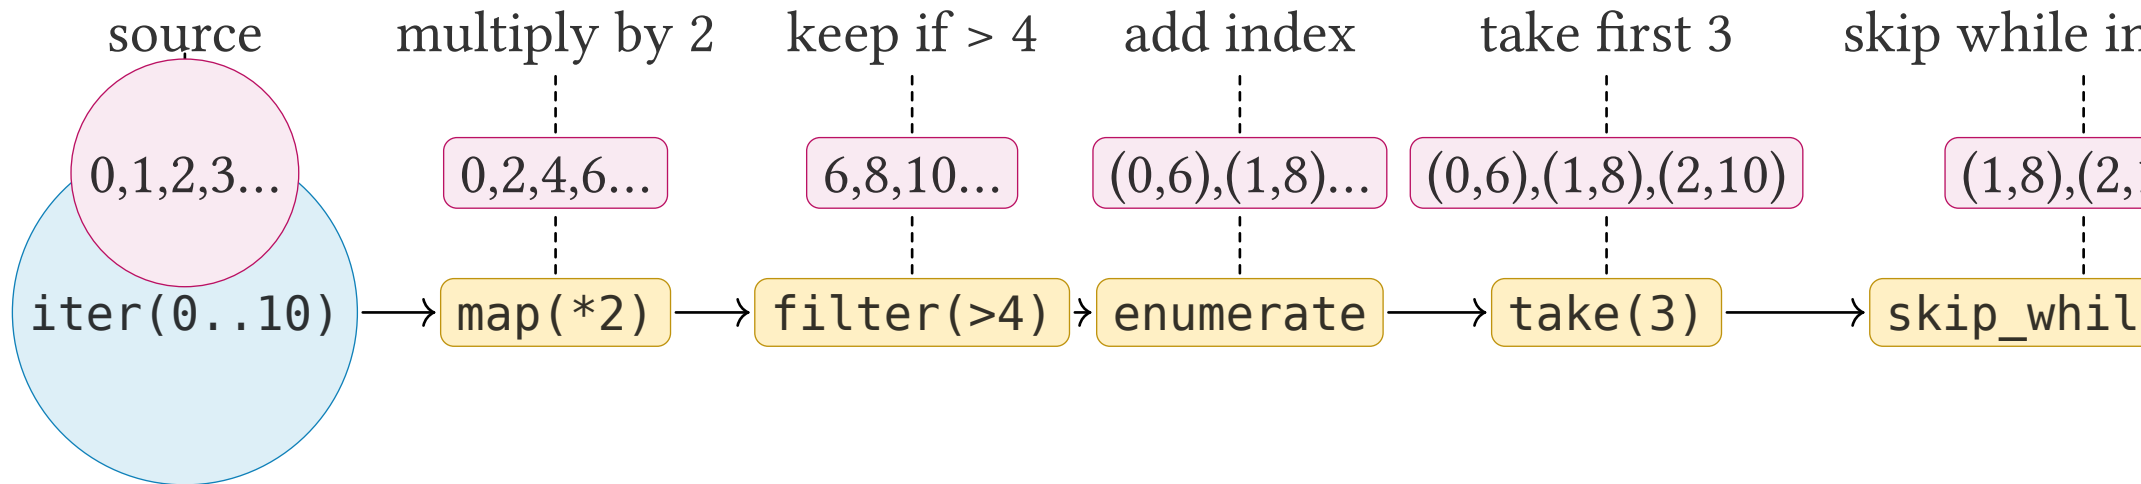


	Future	Stream	Meaning
Regular			May continue
Fused	FusedFuture	FusedStream	is_terminated() method
Fused			Done permanently
Fused value	Pending	Ready (None)	Final value

1.	Introduction	1
2.	Rust's <code>Stream</code> trait	8
3.	Using Streams	14
3.1.	Pipelines with <code>futures::StreamExt</code>	15
3.2.	The handy <code>std::future::ready</code> function	16
3.3.	Flatten a finite collection of <code>Streams</code>	17
3.4.	Flattening an infinite stream	18
3.5.	More <code>Stream</code> features to explore	19
4.	Example 1: $1 \rightarrow 1$ Operator	20
5.	Example 2: $1 \rightarrow N$ Operator	44
6.	Conclusion	51


3.1. Pipelines with futures::StreamExt

All basic stream operators are in `futures::StreamExt`



```

1 stream::iter(0..10)
2   .map(|x| x * 2)
3   .filter(|&x| ready(x > 4))
4   .enumerate()
5   .take(3)
6   .skip_while(|&(i, _)| i < 1)
  
```




Rust

3.2. The handy `std::future::ready` function

The `futures::StreamExt::filter` expects an **async closure** (or closure returning `Future`): 


Option 1: Async block (not `Unpin!`)

```
1 stream.filter(|&x| async move {
2     x % 2 == 0
3 })
```

 Rust


Option 2: Async closure (not `Unpin!`)

```
1 stream.filter(async |&x| x % 2
    == 0)
```

 Rust

Option 3 (recommended): Wrap sync output with `std::future::ready()`

```
1 stream.filter(|&x| ready(x %
    2 == 0))
```


 Rust

- `ready(value)` creates a `Future` that immediately resolves to `value`.
- `ready(value)` is `Unpin`

`ready` keeps pipelines `Unpin`: *easier to work with*

3.3. Flatten a finite collection of Streams

A finite collection of Streams = `IntoIterator<Item: Stream>`



```
1 let streams = vec![  
2     stream::iter(1..=3),  
3     stream::iter(4..=6),  
4     stream::iter(7..=9),  
5 ];  
6  
7 let merged = stream::select_all(streams);
```


1. Creates a `FuturesUnordered` of the streams
2. Polls all streams concurrently
3. Yields items as they arrive

3.4. Flattening an infinite stream

Beware!: `flatten()` on a stream of infinite streams will never complete!




```
1 let infinite_streams = stream::unfold(0, |id| async move {  
2     Some((stream::iter(id..), id + 1))  
3 });  
4 let flat = infinite_streams.flatten();
```



Instead, **buffer streams** concurrently with `flatten_unordered()`.


```
1 let requests = stream::unfold(0, |id| async move {  
2     Some((fetch_stream(format!("/api/data/{}", id)), id + 1))  
3 });  
4 let flat = requests.flatten_unordered(Some(10));
```



3.5. More Stream features to explore

Many more advanced topics await:

- **Boolean operations:** any, all
- **Async operations:** then
- **Sinks:** The write-side counterpart to Streams




1.	Introduction	1
2.	Rust's <code>Stream</code> trait	8
3.	Using Streams	14
4.	Example 1: $1 \rightarrow 1$ Operator	20
4.1.	Concept: Doubling stream operator	22
4.2.	Building a stream operator: structure	23
4.3.	Implementing the <code>Stream</code> trait	24
4.4.	Naive implementation of <code>poll_next</code>	25
4.5.	Accessing pinned fields	26
4.6.	The naive solution fails	27
4.7.	Why does <code>Pin::get_mut()</code> require <code>Unpin</code> ? ...	28
4.8.	<code>Unpin</code> types can be safely unpinned	29
4.9.	More <code>Unpin</code> types	30
4.10.	<code>!Unpin</code> types cannot be safely unpinned	31

4.11.	More !Unpin types	32
4.12.	One workaround: add the Unpin bound	33
4.13.	Turning !Unpin into Unpin with boxing	35
4.14.	Applying the solution: Pin<Box<InSt>>	36
4.15.	Projecting visually	37
4.16.	Complete boxed Stream implementation	38
4.17.	Two ways to handle !Unpin fields	39
4.18.	Approach 3: Projection with pin-project	40
4.19.	Distributing your operator	41
4.20.	The ‘ <i>real</i> ’ stream drivers	42
4.21.	The Stream trait: a lazy query interface	43
5.	Example 2: $1 \rightarrow N$ Operator	44
6.	Conclusion	51



Very simple Stream operator that **doubles every item** in an input stream:




Input stream **needs to yield integers**.

4.2. Building a stream operator: structure

4. Example 1: $1 \rightarrow 1$ Operator

Step 1: Define a struct that wraps the input stream


```
1 struct Double<InSt> {  
2     in_stream: InSt,  
3 }
```



- Generic over stream type (works with any backend)
- Stores input stream by value

4.3. Implementing the Stream trait

Step 2: Implement Stream trait with bounds



```
1  impl<InSt> Stream for Double<InSt>
2  where
3      InSt: Stream<Item = i32>
4  {
5      type Item = i32;
6
7      fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)
8          -> Poll<Option<Self::Item>> {
9          // ... implementation goes here
10     }
11 }
```

4.4. Naive implementation of poll_next

4. Example 1: $1 \rightarrow 1$ Operator

Focus on the implementation of the poll_next method



(Remember that Self = Double<InSt> with field in_stream: InSt):

```
1 fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)  
2     -> Poll<Option<Self::Item>> {  
3     // Cannot access self.in_stream!  
4     Pin::new(&mut self.in_stream) // Not possible!  
5     .poll_next(cx)  
6     .map(|x| x * 2)  
7 }
```

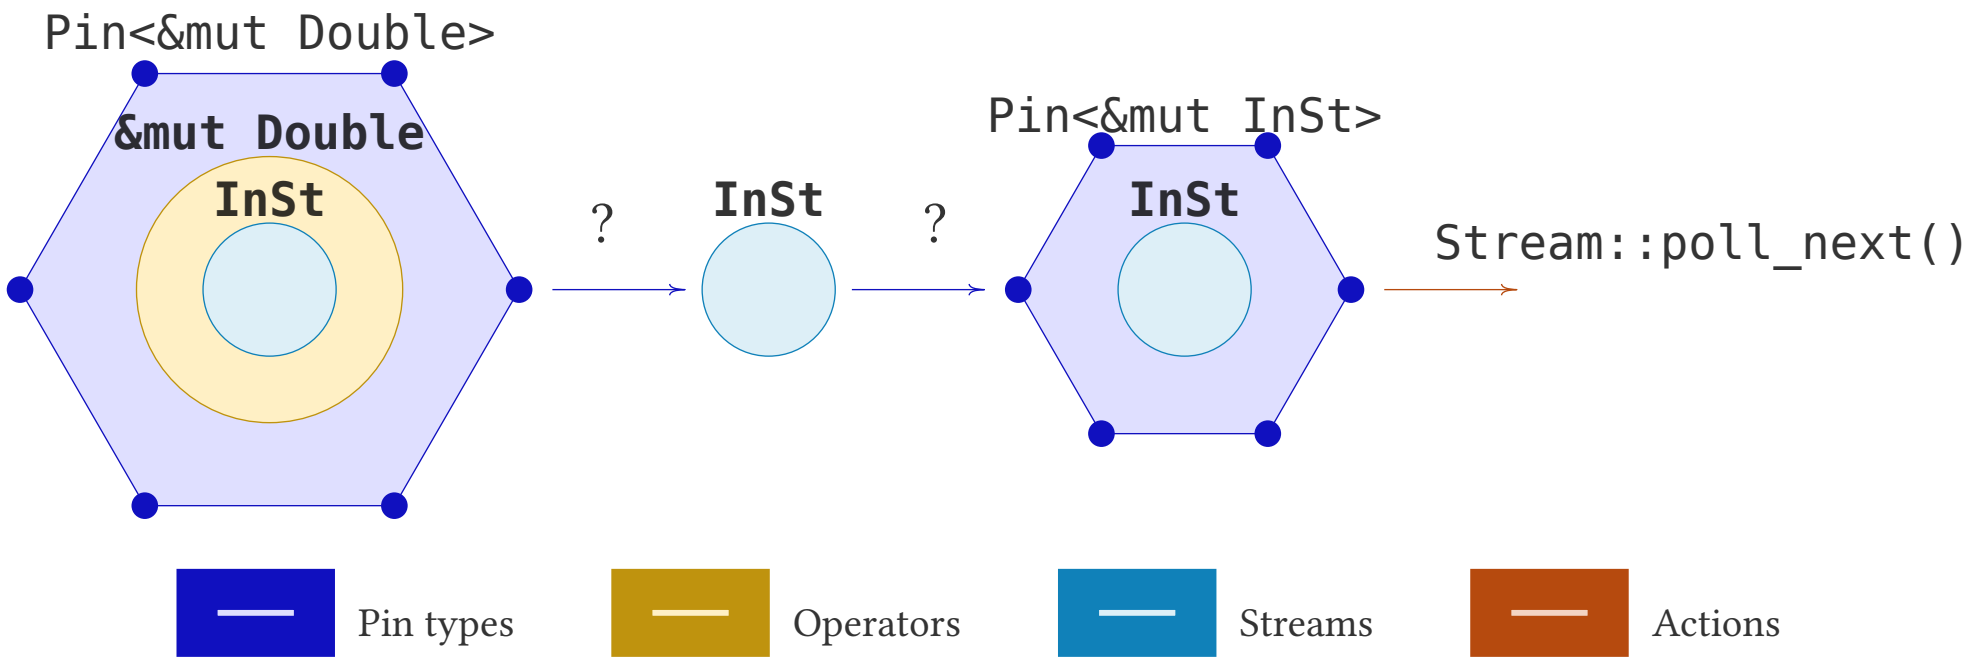


Pin<&mut Self> **blocks access to self.in_stream** (when Self: !Unpin)!

Warning

We have `Pin<&mut Double>`.

How can we obtain `Pin<&mut InSt>` to call `poll_next()`?



4.6. The naive solution fails

Can we use `Pin::get_mut()` to unwrap and re-wrap?

```

1  impl<InSt> Stream for Double<InSt> where InSt: Stream<Item = i32> {
2
3      type Item = InSt::Item;
4
5      fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)
6          -> Poll<Option<Self::Item>> {
7          let this = self.get_mut(); // Error!
8          let pinned_in = Pin::new(&mut this.in_stream);
9          pinned_in.poll_next(cx).map(|p| p.map(|x| x * 2))
10     }
11     }

```



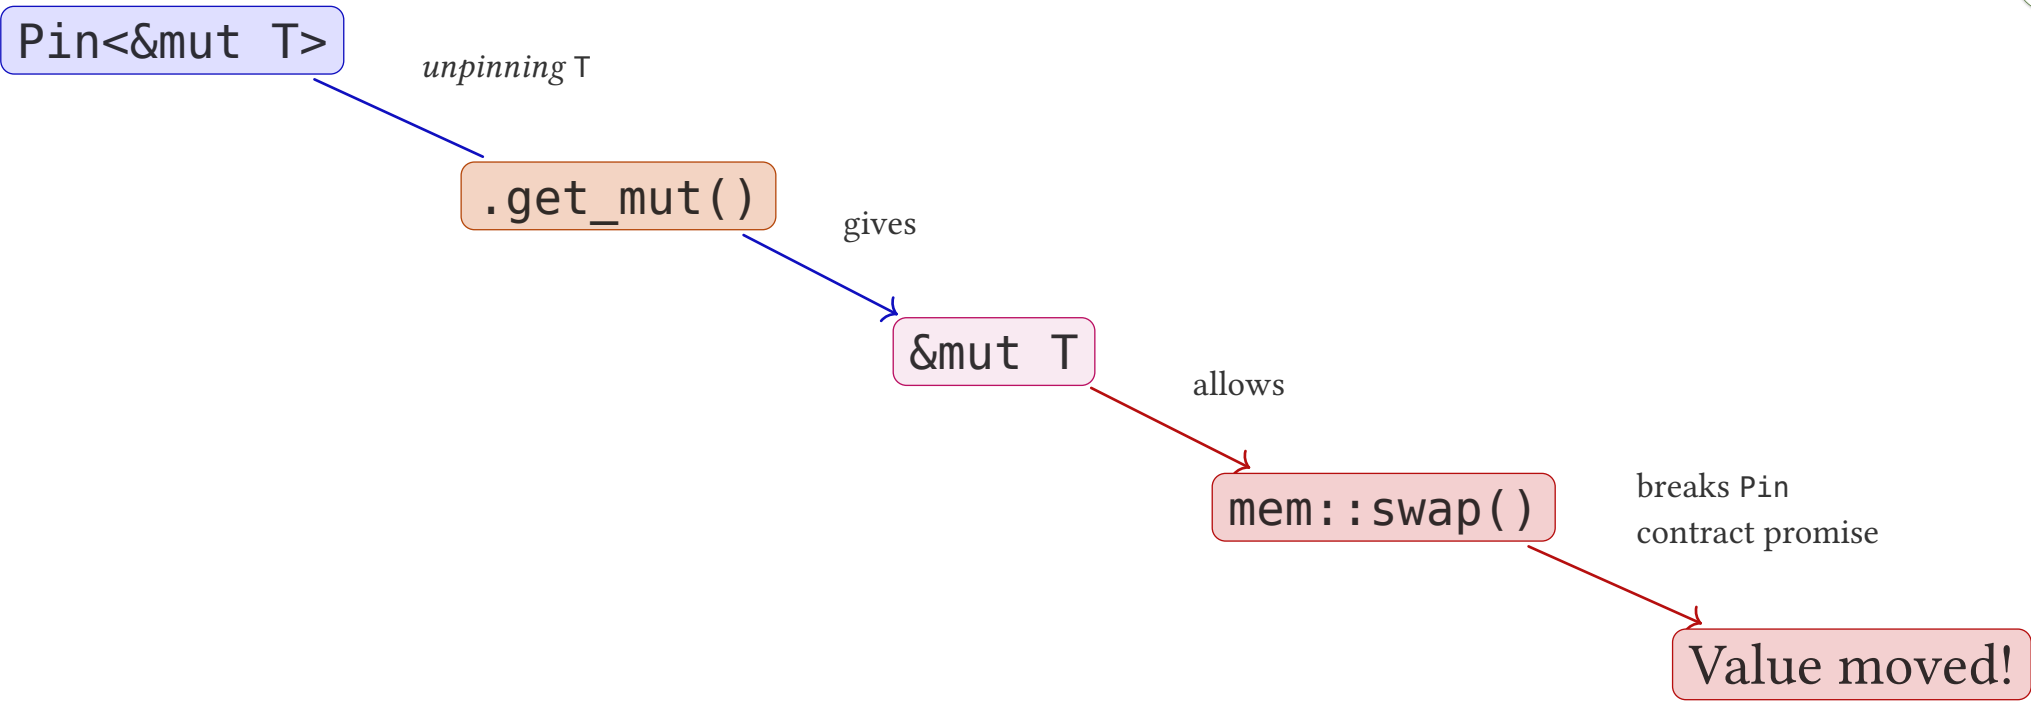
Problem: `Pin::get_mut()` requires `Double<InSt>: Unpin`

But ***Double<InSt> is !Unpin when InSt: !Unpin!***

4.7. Why does `Pin::get_mut()` require `Unpin`?

4. Example 1: $1 \rightarrow 1$ Operator

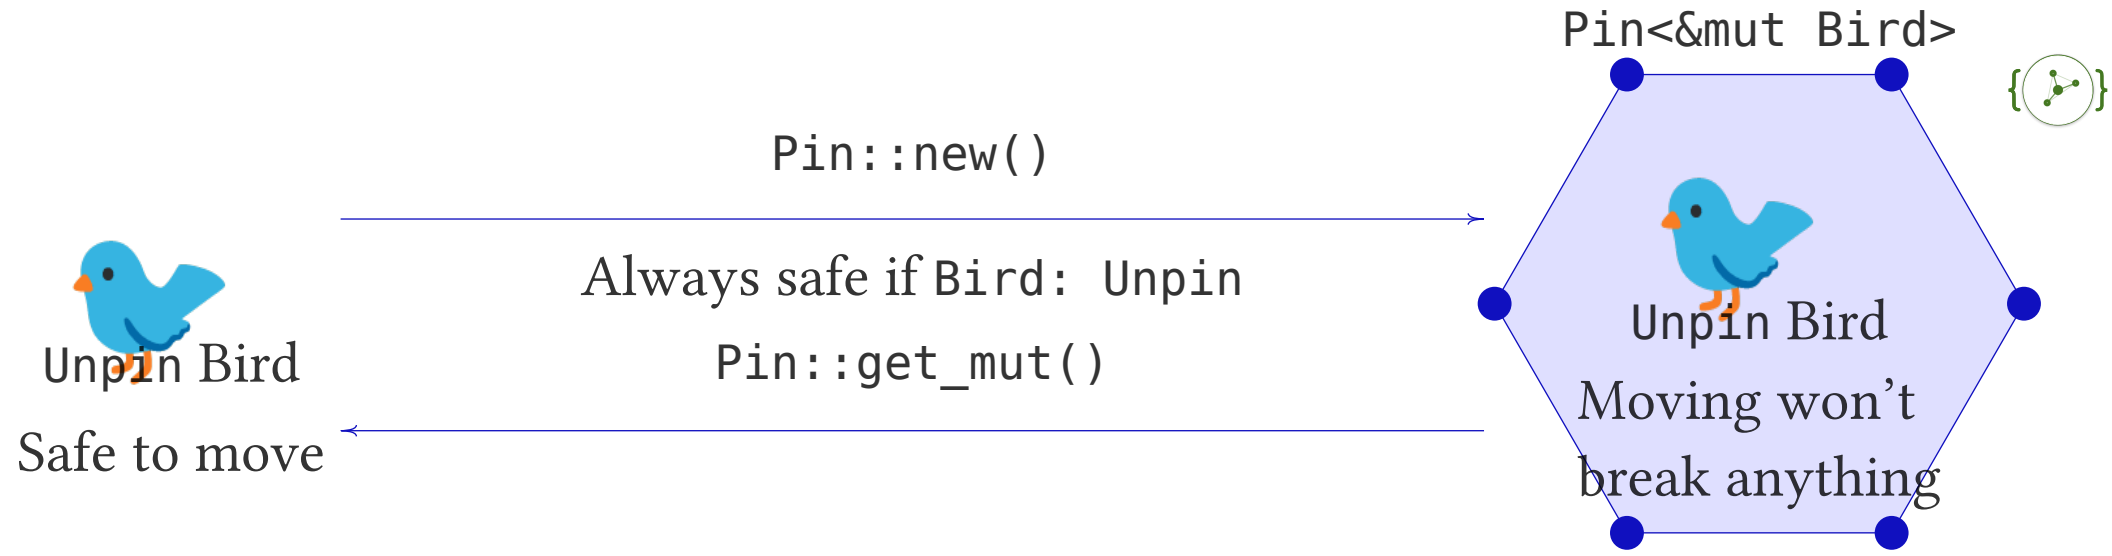
`Pin<P>` makes a promise: **the pointee will never move again.**



Solution: Only allow `get_mut()` when `T: Unpin` (moving is safe).

4.8. Unpin types can be safely unpinned

4. Example 1: $1 \rightarrow 1$ Operator



If `T: Unpin`, then `Pin::get_mut()` is safe because moving `T` doesn't cause UB.

4.9. More Unpin types

Examples of Unpin types:

- `i32`, `String`, `Vec<T>` - all primitive and standard types
- `Box<T>` - pointers are safe to move
- `&T`, `&mut T` - references are safe to move



Why safe?

These types don't have self-referential pointers. Moving them in memory doesn't invalidate any internal references.

Almost all types are `Unpin` by default!

4.10. !Unpin types cannot be safely unpinned

4. Example 1: $1 \rightarrow 1$ Operator




!Unpin Tiger

Dangerous to move

~~Pin::get_mut()
gives &mut T~~

Would break
pin promise!



4.11. More !Unpin types

Examples of !Unpin types:

- PhantomPinned - explicitly opts out of Unpin
- Most Future types (self-ref. state machines)
- Types with self-referential pointers
- Double<InSt> where InSt: !Unpin



Why unsafe?

These types may contain pointers to their own fields. Moving them in memory would invalidate those internal pointers, causing use-after-free.

!Unpin is rare and usually intentional for async/self-referential types.

4.12. One workaround: add the Unpin bound

4. Example 1: $1 \rightarrow 1$ Operator

The compiler error suggests adding `InSt: Unpin`:

```
1  impl<InSt> Stream for Double<InSt> where InSt: Stream<Item = i32> + Unpin {
2      type Item = InSt::Item;
3
4      fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Option<Self::Item>> {
5          // `this` = a conventional name for `get_mut` output
6          let mut this = self.get_mut();
7          let pinned_in = Pin::new(&mut this.in_stream);
8          pinned_in
9              .poll_next(cx)
10             .map(|p| p.map(|x| x * 2))
11     }
12 }
```



Warning

We **don't** want to impose `InSt: Unpin` on users of `Double`!

4.12. One workaround: add the Unpin bound

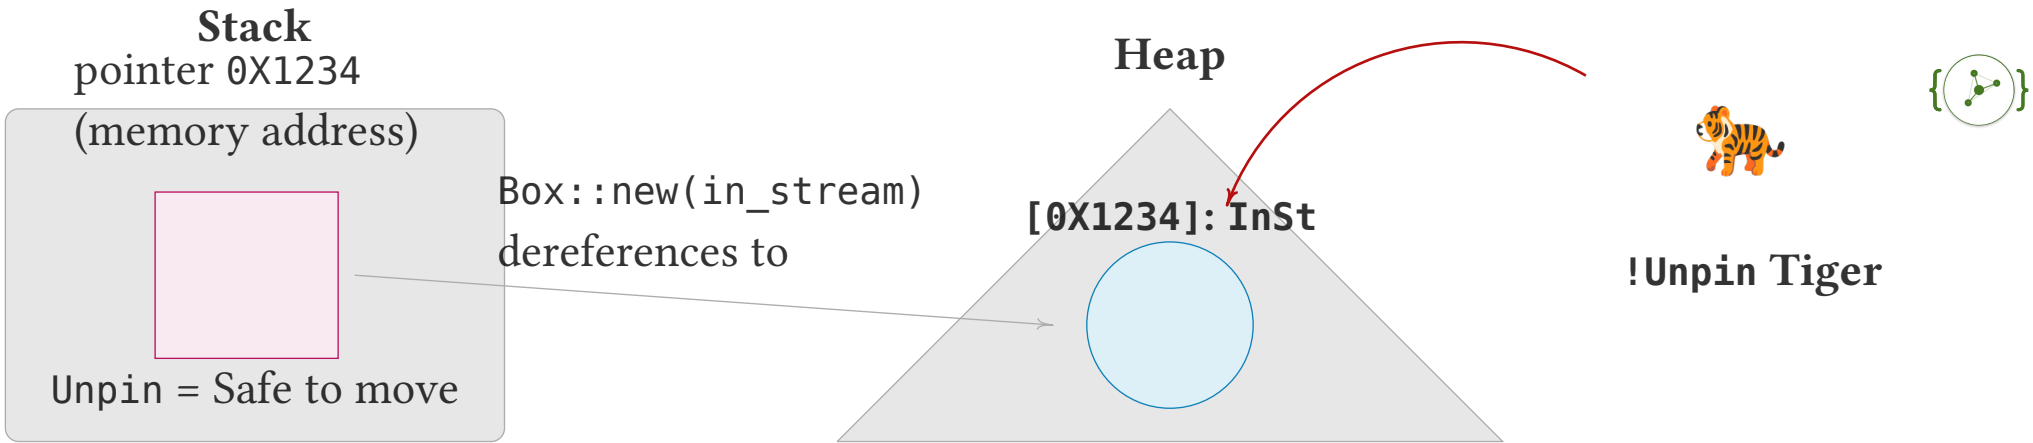
How to support InSt: !Unpin streams? ...

4. Example 1: $1 \rightarrow 1$ Operator



4.13. Turning !Unpin into Unpin with boxing

4. Example 1: $1 \rightarrow 1$ Operator



Nice to have:

- 1. `Box::new(tiger)` produces just a pointer on the stack
 - Moving pointers is always safe
 - Therefore: **`Box<Tiger>: Unpin`**
- 2. Box dereferences to its contents
 - **`Box<X>: Deref<Target = X>`**


Problem: Need `Pin<&mut InSt>`, but `Box<InSt>` requires `InSt: Unpin` to create it

Solution: Use `Pin<Box<InSt>>` to project from `Pin<&mut Double>` to `Pin<&mut InSt>` via `Pin::as_mut()`

4.14. Applying the solution: `Pin<Box<InSt>>`

Change the struct definition to store `Pin<Box<InSt>>`:

```
1 struct Double<InSt> { in_stream: Pin<Box<InSt>>, }
```



Why this works:


- `Box<InSt>` is always `Unpin` (pointers are safe to move)
- `Pin<Box<InSt>>` can hold `!Unpin` streams safely on the heap

Projection in `poll_next`:

```

1 fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)
2     -> Poll<Option<Self::Item>> {
3     let this = self.get_mut(); // Safe: Double is Unpin now
4     this.in_stream.as_mut()     // Project to Pin<&mut InSt>
5         .poll_next(cx)
6         .map(|opt| opt.map(|x| x * 2))
7     }


```



This works **without requiring `InSt: Unpin`**!

4.15. Projecting visually

From `Pin<&mut Double>` to `Pin<&mut InSt>` in a few **safe steps**:



4.16. Complete boxed Stream implementation

4. Example 1: $1 \rightarrow 1$ Operator

We can call `Pin::get_mut()` to get `&mut Double<InSt>` safely from `Pin<&mut Double<InSt>>`


```
1  impl<InSt> Stream for Double<InSt>
2  where InSt: Stream<Item = i32>
3  {
4      fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)
5          -> Poll<Option<Self::Item>>
6      {
7          // We can project because `Self: Unpin`
8          let this: &mut Double<InSt> = self.get_mut();
9          this.in_stream.as_mut()
10             .poll_next(cx)
11             .map(|r| r.map(|x| x * 2))
12      }
13 }
```



4.17. Two ways to handle !Unpin fields

4. Example 1: $1 \rightarrow 1$ Operator


Approach 1: Use Box<_>

```
1 struct Double<InSt> {  Rust
2   in_stream: Pin<Box<InSt>>
3 }
4
5 impl<InSt> Stream for
  Double<InSt>
6   where InSt: Stream
```

✓ Works with any InSt, also !Unpin

... or, use pin-project crate

Approach 2: Require Unpin

```
1 struct Double<InSt> {  Rust
2   in_stream: InSt
3 }
4
5 impl<InSt> Stream for
  Double<InSt>
6   where InSt: Stream + Unpin
```

✗ Imposes Unpin constraint on users



4.18. Approach 3: Projection with pin-project

Projects like Tokio use the pin-project crate:



```
1  #[pin_project]
2  struct Double<InSt> {
3      #[pin]
4      in_stream: InSt,
5  }
6  impl<InSt: Stream> Stream for Double<InSt> {
7      fn poll_next(self: Pin<&mut Self>, cx: &mut Context<'_>)
8          -> Poll<Option<Self::Item>>
9      {
10         // pin-project generates a safe projection method `project()`
11         self.project().in_stream.poll_next(cx)
12             .map(|r| r.map(|x| x * 2))
13     }
14 }
```

Uses a macros underneath.

4.19. Distributing your operator

Define a constructor and turn it into a method of an **extension trait**:



```
1 trait DoubleStream: Stream {  
2     fn double(self) -> Double<Self>  
3     where Self: Sized + Stream<Item = i32>,  
4     { Double::new(self) }  
5 }  
6 // A blanket implementation should be provided by you!  
7 impl<S> DoubleStream for S where S: Stream<Item = i32> {}
```

Now, users **don't need to know how** Double is implemented, just

1. import your extension trait: DoubleStream
2. call `.double()` on any compatible stream



Stream Trait Interface

Lazy: `.poll_next()` only responds when called



Data pushed up

Leaf Streams (Real Drivers)

TCP, Files, Timers, Hardware, Channels

Stream trait just provides a **uniform way to query** - it doesn't create or drive data flow.

4.21. The Stream trait: a lazy query interface

4. Example 1: $1 \rightarrow 1$ Operator

The **Stream** trait is **NOT** the stream itself - it's just a lazy frontend to query data.



What **Stream** trait does:

- Provides uniform `.poll_next()` interface
- Lazy: only responds when asked
- Doesn't drive or produce data itself
- Just queries whatever backend exists

What actually drives streams:

- TCP connections receiving packets
- File I/O completing reads
- Timers firing
- Hardware signals
- Channel senders pushing data

1.	Introduction	1
2.	Rust's Stream trait	8
3.	Using Streams	14
4.	Example 1: $1 \rightarrow 1$ Operator	20
5.	Example 2: $1 \rightarrow N$ Operator	44
5.1.	Complexity $1 \rightarrow N$ operators	45
5.2.	Sharing latency between tasks	46
5.3.	Cloning streams with an operator	47
5.4.	Polling and waking flow	48
5.5.	Steps for creating robust stream operators	49
5.6.	Simplified state machine of clone-stream	50
6.	Conclusion	51

5.1. Complexity $1 \rightarrow N$ operators

5. Example 2: $1 \rightarrow N$ Operator


Challenges for Stream operators are combined from:

Inherent Future challenges:

- Clean up orphaned wakers
- Cleanup when tasks abort
- Task coordination complexity

Inherent Iterator challenges:

- Ordering guarantees across consumers
- Backpressure with slow consumers
- Sharing mutable state safely
- Avoiding duplicate items



5.2. Sharing latency between tasks

5. Example 2: $1 \rightarrow N$ Operator

Latency may need to be processed by different async tasks:



```
1 let tcp_stream =  
  TcpStream::connect("127.0.0.1:8080").await?;  
2 let latency = tcp_stream.latency(); // Stream<Item = Duration>  
3  
4 spawn(async move { display_ui(latency).await; });  
5 spawn(async move { engage_breaks(latency).await; }); // Error!
```



Error

latency is moved into the first task, so the second task can't access it.

Warning

We need a way to clone the latency stream!


5.3. Cloning streams with an operator

5. Example 2: $1 \rightarrow N$ Operator

Solution: Create a *stream operator* `fork()` makes the input stream Clone.



```
1 let ui_latency = tcp_stream.latency().fork();
```



```
2
```

```
3 let breaks_latency_clone = ui_latency.clone();
```

```
4 // Warning: `Clone` needs to be implemented!
```


```
5
```

```
6 spawn(async move { display_ui(ui_latency).await; });
```

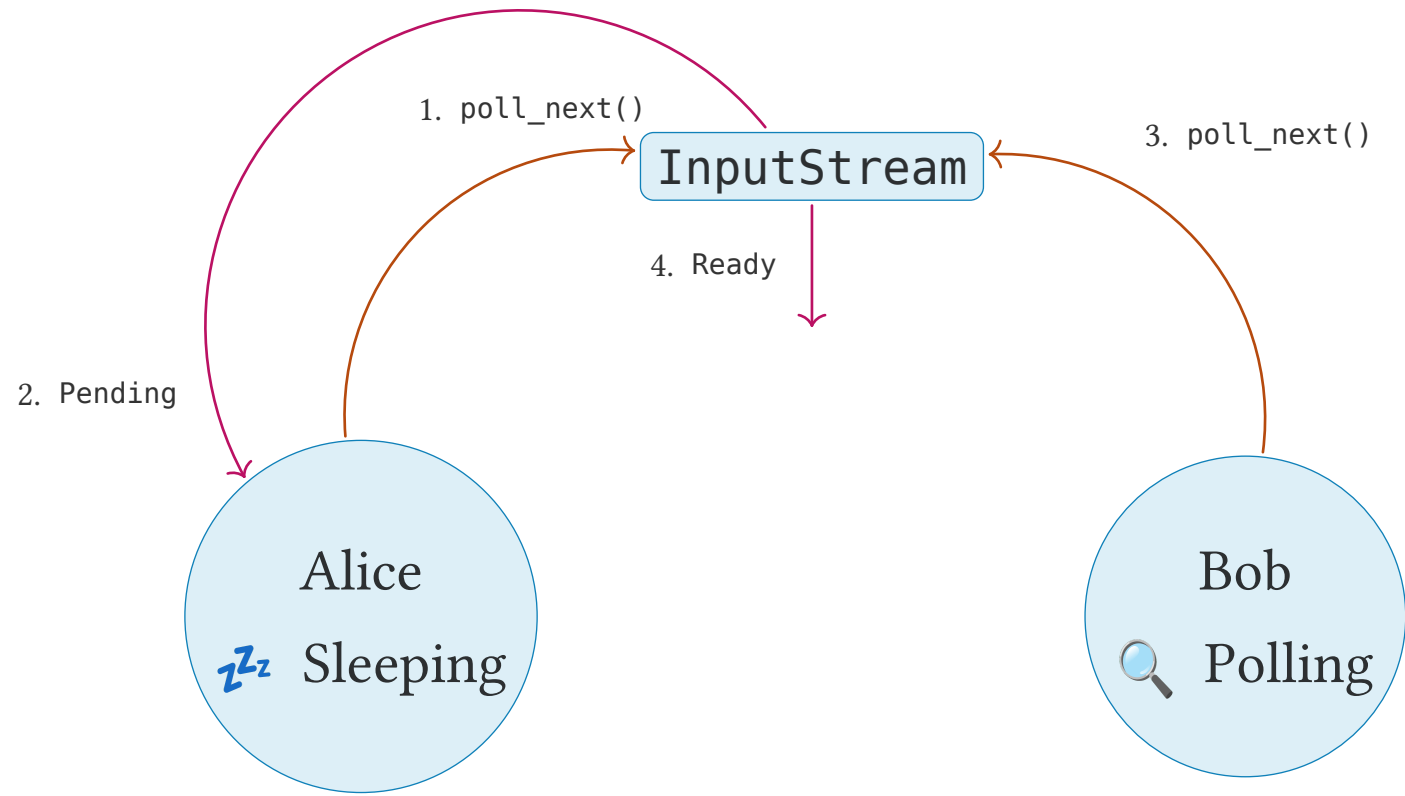
```
7 spawn(async move  
  { engage_breaks(breaks_latency_clone).await; });
```

Requirement: `Stream<Item: Clone>`, so we can clone the items (Duration is Clone)

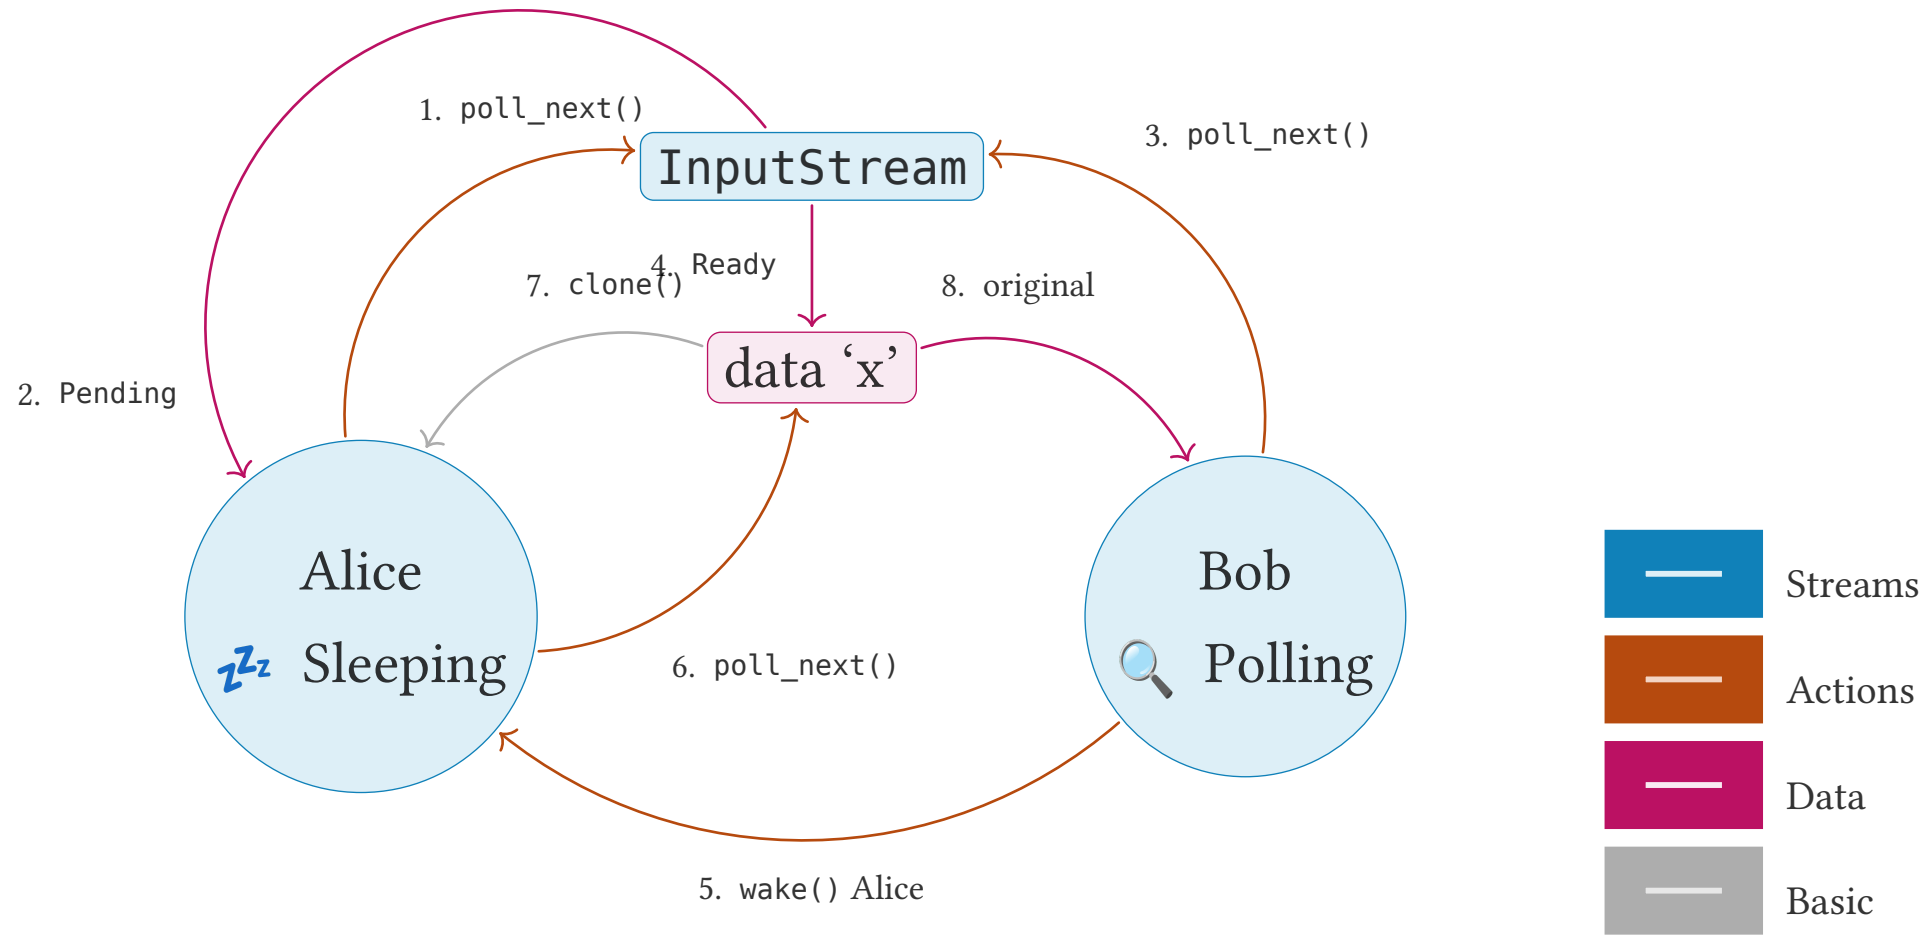
5.4. Polling and waking flow



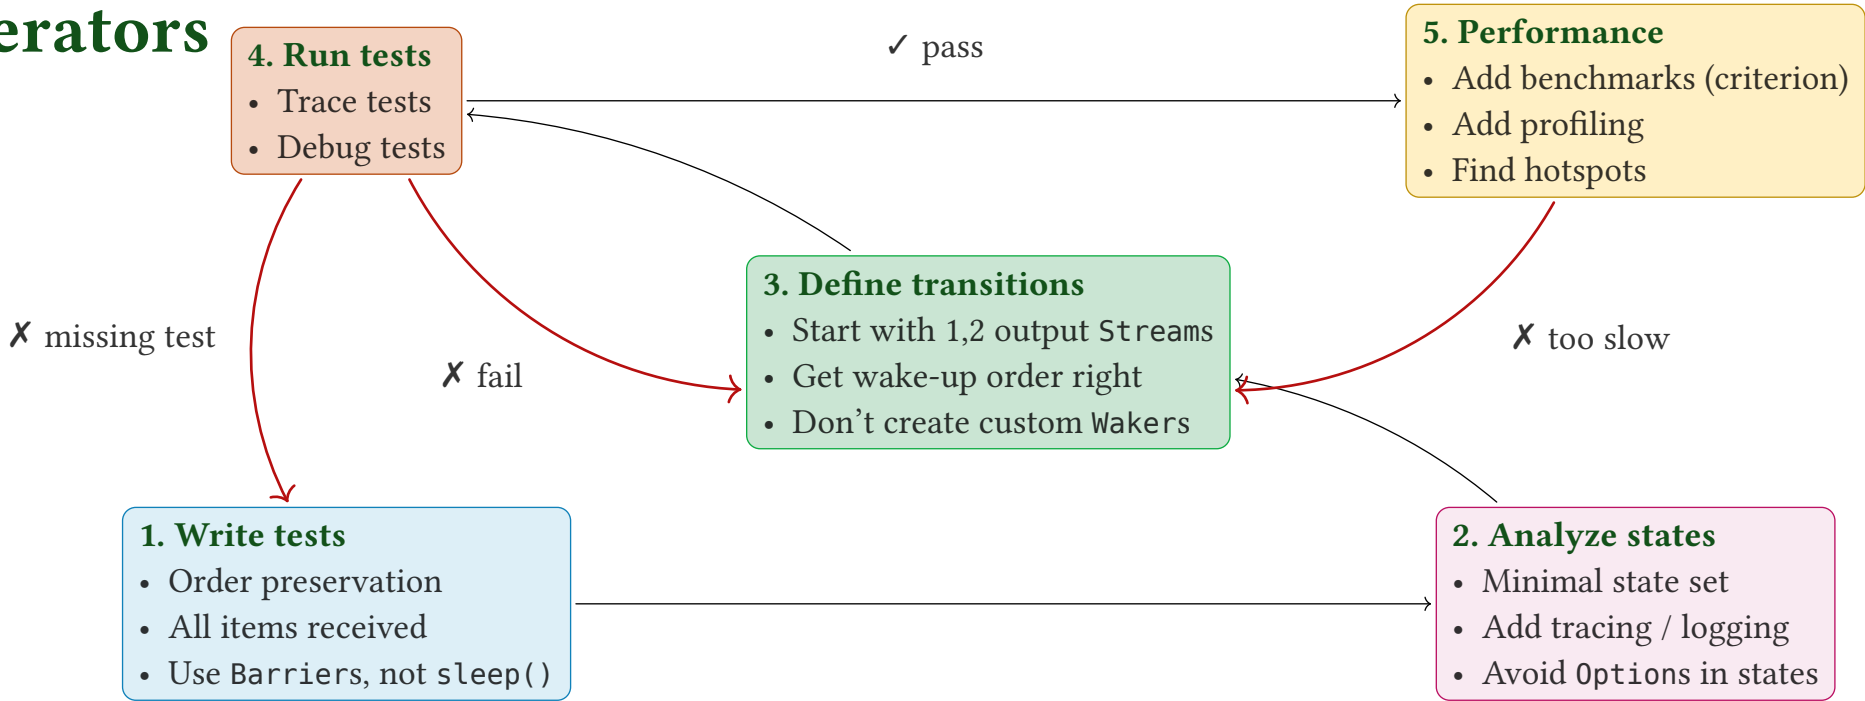
5.4. Polling and waking flow



5.4. Polling and waking flow



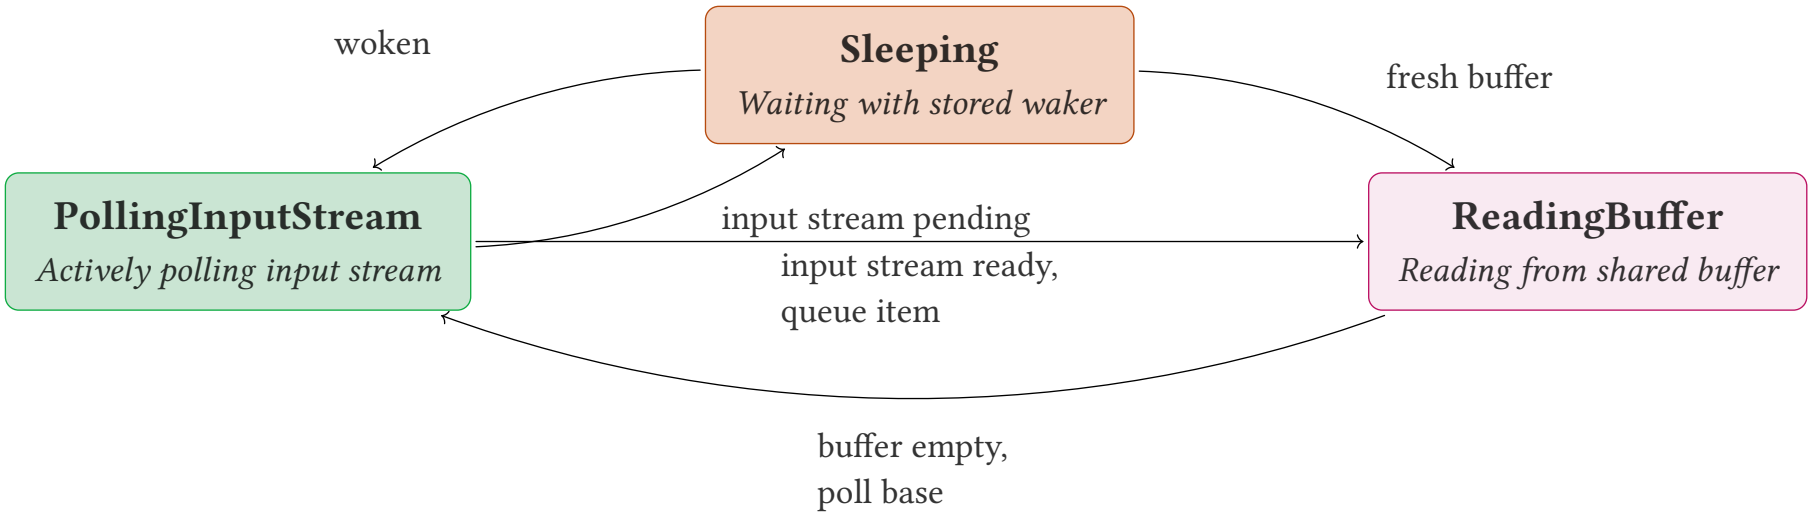
5.5. Steps for creating robust stream operators



5.6. Simplified state machine of clone-stream

5. Example 2: 1 → N Operator



Enforcing simplicity, correctness and performance:



Warning

Each clone maintains its own state

1.	Introduction	1
2.	Rust's Stream trait	8
3.	Using Streams	14
4.	Example 1: $1 \rightarrow 1$ Operator	20
5.	Example 2: $1 \rightarrow N$ Operator	44
6.	Conclusion	51
6.1.	Quickstart	52
6.2.	Questions	53



6.2. Questions

Thank you for your attention!

- Contact me: willemvanhulle@protonmail.com
- These slides: github.com/wvhulle/streams-eurorust-2025



Want to learn more?

Join my 7-week course ***“Creating Safe Systems in Rust”***

- Location: Ghent (Belgium)
- Date: starting 4th of November 2025.

Register at pretix.eu/devlab/rust-course/